


```

17 0.14          11429 non-null float64
18 0.15          11429 non-null float64
19 0.5.1         11429 non-null float64
20 0.16          11429 non-null float64
21 0.17          11429 non-null int64
22 0.18          11429 non-null float64
23 1             11429 non-null int64
24 0.19          11429 non-null int64
25 0.20          11429 non-null int64
26 0.21          11429 non-null int64
27 0.22          11429 non-null int64
28 0.23          11429 non-null int64
29 0.24          11429 non-null int64
30 1.1           11429 non-null float64
31 0.027397260273972601 11429 non-null float64
32 0.25          11429 non-null int64
33 0.26          11429 non-null int64
dtypes: float64(20), int64(14)
memory usage: 3.0 MB

```

```

In [46]: from sklearn.model_selection import train_test_split

arr = data.values
X = arr[:,0:33]
Y = arr[:,33]

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.30)

```

```

In [47]: results = {
    "KNN": 0,
    "SVM": 0,
    "Naive Bayes": 0,
    "Decision Tree": 0,
}

cv = KFold(n_splits=10, random_state=1, shuffle=True)

from sklearn.metrics import classification_report, confusion_matrix

def report(y_test, y_pred, classifier):
    scores = cross_val_score(classifier, X, Y, scoring='accuracy', cv=cv, n_jobs=
    avg_accuracy = sum(scores) / len(scores)

    return avg_accuracy

```

KNN Classifier

```

In [48]: from sklearn.neighbors import KNeighborsClassifier

scores = []

for n in range(1, 25, 2):
    knn_classifier = KNeighborsClassifier(n_neighbors=n)

    knn_classifier.fit(X_train, y_train)

```

```
y_pred = knn_classifier.predict(X_test)

accuracy = report(y_test, y_pred, knn_classifier)

scores.append((n, accuracy))
```

In [49]: scores

```
Out[49]: [(1, 0.8212444438315613),
          (3, 0.833582010654973),
          (5, 0.8262327760693662),
          (7, 0.825445374494563),
          (9, 0.8181830927001025),
          (11, 0.8153838257083015),
          (13, 0.8168716760667614),
          (15, 0.8126722776115333),
          (17, 0.8103969490678814),
          (19, 0.8054097659859067),
          (21, 0.8053219704804851),
          (23, 0.8050589670161633)]
```

```
In [50]: max_performace = max(scores, key=lambda item:item[1])

best_k = max_performace[0]
score = max_performace[1]

results["KNN"] = score


print ("Best K:", best_k)
print ("KNN Accuracy Score:", score)
```

Best K: 3
KNN Accuracy Score: 0.833582010654973

```
In [51]: x = [k for k, accuracy in scores]
y = [accuracy for k, accuracy in scores]

colors = np.random.rand(len(y))

plt.scatter(x, y, c=colors, alpha=1)
plt.xlabel('K')
plt.ylabel('Accuracy')
plt.plot(x, y)
plt.show()
```



SVM

```
In [52]: from sklearn import svm

kernels = ["poly", "rbf", "sigmoid", "linear"]
scores = []

for kernel in kernels:
    clf = svm.SVC(kernel=kernel)

    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)

    accuracy = report(y_test, y_pred, clf)

    scores.append((kernel, accuracy))
```

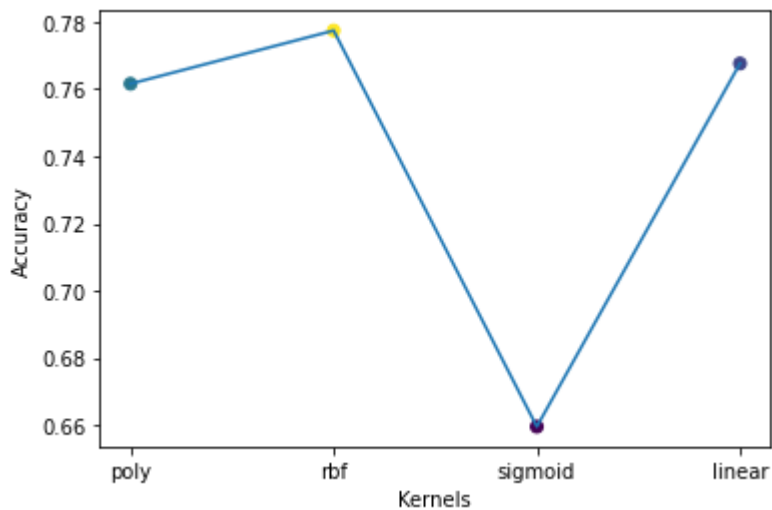
```
In [53]: scores
```

```
Out[53]: [('poly', 0.7615748337937618),
          ('rbf', 0.7774086689251409),
          ('sigmoid', 0.6595510171561304),
          ('linear', 0.7676089744473711)]
```

```
In [54]: x = [kernel for kernel, accuracy in scores]
y = [accuracy for kernel, accuracy in scores]

colors = np.random.rand(len(y))

plt.scatter(x, y, c=colors, alpha=1)
plt.xlabel('Kernels')
plt.ylabel('Accuracy')
plt.plot(x, y)
plt.show()
```



```
In [55]: max_performace = max(scores,key=lambda item:item[1])

best_kernel = max_performace[0]
score = max_performace[1]

results["SVM"] = score

print ("Best Kernel :", best_kernel)
print ("SVM Accuracy Score :", score)
```

Best Kernel : rbf
SVM Accuracy Score : 0.7774086689251409

Naive Bayes

```
In [56]: from sklearn.naive_bayes import GaussianNB

gnb = GaussianNB()

gnb.fit(X_train, y_train)

y_pred = gnb.predict(X_test)

accuracy = report(y_test, y_pred, gnb)

results["Naive Bayes"] = accuracy

print ("Avg. accuracy ", accuracy)

print (classification_report(y_test, y_pred))
```

Avg. accuracy	0.617988885364811				
	precision	recall	f1-score	support	
0.0	0.58	0.97	0.72	1721	
1.0	0.91	0.28	0.42	1708	
accuracy			0.63	3429	
macro avg	0.74	0.62	0.57	3429	
weighted avg	0.74	0.63	0.57	3429	

Decision Tree

```
In [57]: from sklearn.tree import DecisionTreeClassifier

DTC = DecisionTreeClassifier()

DTC.fit(X_train,y_train)

y_pred = DTC.predict(X_test)

accuracy = report(y_test, y_pred, DTC)

results["Decision Tree"] = accuracy

print ("Decision Tree accuracy score: ", accuracy)

print (classification_report(y_test, y_pred))
```

```
Decision Tree accuracy score: 0.822819170370779
              precision    recall  f1-score   support

    0.0         0.82        0.83        0.83        1721
    1.0         0.83        0.82        0.83        1708

 accuracy                   0.83        3429
 macro avg                 0.83        0.83        0.83        3429
weighted avg                 0.83        0.83        0.83        3429
```

Logistic Regression

```
In [58]: from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression()

logreg.fit(X_train,y_train)

y_pred = logreg.predict(X_test)

accuracy = report(y_test, y_pred, logreg)

results["Logistic Regression"] = accuracy

print ("Avg. accuracy ", accuracy)

print (classification_report(y_test, y_pred))
```

```
Avg. accuracy 0.7785465630281329
              precision    recall  f1-score   support

    0.0         0.79        0.77        0.78        1721
    1.0         0.77        0.80        0.78        1708

 accuracy                   0.78        3429
 macro avg                 0.78        0.78        0.78        3429
weighted avg                 0.78        0.78        0.78        3429
```

Summarizing Results

```
In [59]: results
```

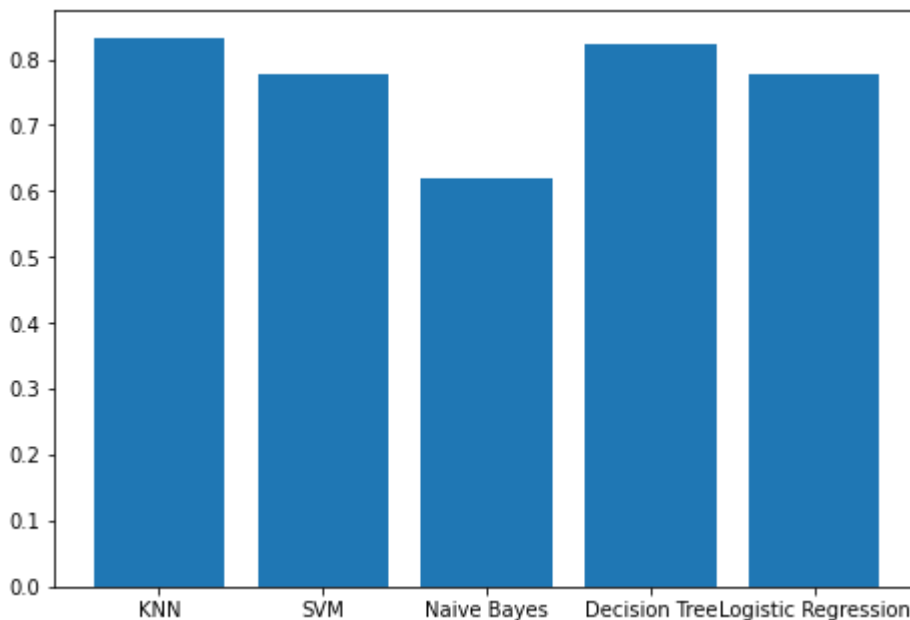
```
Out[59]: {'KNN': 0.833582010654973,  
          'SVM': 0.7774086689251409,  
          'Naive Bayes': 0.617988885364811,  
          'Decision Tree': 0.822819170370779,  
          'Logistic Regression': 0.7785465630281329}
```

```
In [60]: results.keys()  
scores = [results[key] for key in results]
```

```
In [61]: results
```

```
Out[61]: {'KNN': 0.833582010654973,  
          'SVM': 0.7774086689251409,  
          'Naive Bayes': 0.617988885364811,  
          'Decision Tree': 0.822819170370779,  
          'Logistic Regression': 0.7785465630281329}
```

```
In [62]: fig = plt.figure()  
ax = fig.add_axes([1,0,1,1])  
  
models = [model for model in results]  
scores = [results[model] for model in results]  
ax.bar(models,scores)  
plt.show()
```



```
In [63]: from sklearn.ensemble import RandomForestClassifier  
  
#Create a Gaussian Classifier  
clf=RandomForestClassifier(n_estimators=100)  
  
#Train the model using the training sets y_pred=clf.predict(X_test)  
clf.fit(X_train,y_train)  
  
y_pred=clf.predict(X_test)
```

```

accuracy = report(y_test, y_pred, logreg)

print ("Avg. accuracy ", accuracy)

print (classification_report(y_test, y_pred))

```

```

Avg. accuracy 0.7785465630281329
              precision    recall  f1-score   support

    0.0         0.86         0.85         0.85         1721
    1.0         0.85         0.86         0.86         1708

 accuracy          0.86          0.86          0.86          3429
  macro avg         0.86          0.86          0.86          3429
 weighted avg         0.86          0.86          0.86          3429

```

In [64]: `print("Accuracy:", metrics.accuracy_score(y_test, y_pred))`

```

Accuracy: 0.8556430446194225

```